

Problem statement: Performing exploratory data analysis using iris dataset

```
#import the useful libraries.
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

# Read the data set of "Iris" in data.
df= pd.read_csv("Iris dataset.csv")

# Printing the data
df
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa
...
145	146	6.7	3.0	5.2	2.3	Iris-virginica
146	147	6.3	2.5	5.0	1.9	Iris-virginica
147	148	6.5	3.0	5.2	2.0	Iris-virginica
148	149	6.2	3.4	5.4	2.3	Iris-virginica
149	150	5.9	3.0	5.1	1.8	Iris-virginica

150 rows × 6 columns

```
#print the head of the data frame.
df.head()
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa

```
df.shape
```

```
(150, 6)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Id               150 non-null    int64
1   SepalLengthCm   150 non-null    float64
2   SepalWidthCm    150 non-null    float64
3   PetalLengthCm   150 non-null    float64
4   PetalWidthCm    150 non-null    float64
5   Species         150 non-null    object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
```

```
df.describe()
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	150.000000	150.000000	150.000000	150.000000

```
df.isnull().sum()
```

```

Id          0
SepalLengthCm  0
SepalWidthCm  0
PetalLengthCm  0
PetalWidthCm  0
Species      0
dtype: int64

```

```
data = df.drop_duplicates(subset ="Species",)
data
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
50	51	7.0	3.2	4.7	1.4	Iris-versicolor
100	101	6.3	3.3	6.0	2.5	Iris-virginica

```
df.value_counts("Species")
```

```

Species
Iris-setosa      50
Iris-versicolor  50
Iris-virginica   50
dtype: int64

```

```

# importing packages
import seaborn as sns
import matplotlib.pyplot as plt

```

```

sns.countplot(x='Species', data=df, )
plt.show()

```

```

sns.scatterplot(x='SepalLengthCm', y='SepalWidthCm',
                hue='Species', data=df, )

```

```
# Placing Legend outside the Figure
```

```
plt.legend(bbox_to_anchor=(1, 1), loc=2)
```

```
plt.show()
```

```
# importing packages
import seaborn as sns
import matplotlib.pyplot as plt
```

```
sns.scatterplot(x='PetalLengthCm', y='PetalWidthCm',
                hue='Species', data=df, )
```

```
# Placing Legend outside the Figure
plt.legend(bbox_to_anchor=(1, 1), loc=2)
```

```
plt.show()
```

```
# importing packages
import seaborn as sns
import matplotlib.pyplot as plt
```

```
sns.pairplot(df.drop(['Id'], axis = 1),
             hue='Species', height=2)
```

```
# importing packages
import seaborn as sns
import matplotlib.pyplot as plt
```

```
fig, axes = plt.subplots(2, 2, figsize=(10,10))
```

```
axes[0,0].set_title("Sepal Length")
axes[0,0].hist(df['SepalLengthCm'], bins=7)
```

```
axes[0,1].set_title("Sepal Width")
axes[0,1].hist(df['SepalWidthCm'], bins=5);
```

```
axes[1,0].set_title("Petal Length")
axes[1,0].hist(df['PetalLengthCm'], bins=6);
```

```
axes[1,1].set_title("Petal Width")
axes[1,1].hist(df['PetalWidthCm'], bins=6);
```

```
# importing packages
import seaborn as sns
import matplotlib.pyplot as plt
```

```
plot = sns.FacetGrid(df, hue="Species")
plot.map(sns.distplot, "SepalLengthCm").add_legend()
```

```
plot = sns.FacetGrid(df, hue="Species")
plot.map(sns.distplot, "SepalWidthCm").add_legend()

plot = sns.FacetGrid(df, hue="Species")
plot.map(sns.distplot, "PetalLengthCm").add_legend()

plot = sns.FacetGrid(df, hue="Species")
plot.map(sns.distplot, "PetalWidthCm").add_legend()

plt.show()
```

```

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning:
warnings.warn(msg, FutureWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning:
warnings.warn(msg, FutureWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning:
warnings.warn(msg, FutureWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning:
warnings.warn(msg, FutureWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning:
warnings.warn(msg, FutureWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning:
warnings.warn(msg, FutureWarning)

```

```
data.corr(method='pearson')
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
Id	1.000000	0.624413	-0.654654	0.969909	0.999685
SepalLengthCm	0.624413	1.000000	-0.999226	0.795795	0.643817
SepalWidthCm	-0.654654	-0.999226	1.000000	-0.818999	-0.673417
PetalLengthCm	0.969909	0.795795	-0.818999	1.000000	0.975713
PetalWidthCm	0.999685	0.643817	-0.673417	0.975713	1.000000

```

# importing packages
import seaborn as sns
import matplotlib.pyplot as plt

```

```

sns.heatmap(df.corr(method='pearson').drop(
    ['Id'], axis=1).drop(['Id'], axis=0),
    annot = True);

```

```
plt.show()
```

```

# importing packages
import seaborn as sns
import matplotlib.pyplot as plt

```

```

def graph(y):
    sns.boxplot(x="Species", y=y, data=df)

```

```
plt.figure(figsize=(10,10))
```

```

# Adding the subplot at the specified
# grid position
plt.subplot(221)
graph('SepalLengthCm')

```


```
plt.subplot(222)
```

```
graph('SepalWidthCm')

plt.subplot(223)
graph('PetalLengthCm')

plt.subplot(224)
graph('PetalWidthCm')

plt.show()
```




```
# importing packages
import seaborn as sns
import matplotlib.pyplot as plt

# Load the dataset
df = pd.read_csv('Iris dataset.csv')

sns.boxplot(x='SepalWidthCm', data=df)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f32698349d0>



```
# Importing
import sklearn
from sklearn.datasets import load_boston
import pandas as pd
import seaborn as sns

# Load the dataset
df = pd.read_csv('Iris dataset.csv')
```

```

# IQR
Q1 = np.percentile(df['SepalWidthCm'], 25,
                    interpolation = 'midpoint')

Q3 = np.percentile(df['SepalWidthCm'], 75,
                    interpolation = 'midpoint')
IQR = Q3 - Q1

print("Old Shape: ", df.shape)

# Upper bound
upper = np.where(df['SepalWidthCm'] >= (Q3+1.5*IQR))

# Lower bound
lower = np.where(df['SepalWidthCm'] <= (Q1-1.5*IQR))

# Removing the Outliers
df.drop(upper[0], inplace = True)
df.drop(lower[0], inplace = True)

print("New Shape: ", df.shape)


sns.boxplot(x='SepalWidthCm', data=df)

```

```

Old Shape: (150, 6)
New Shape: (146, 6)
<matplotlib.axes._subplots.AxesSubplot at
0x7f326977d650>

```



► REGRESSION MODEL

[] ↳ 5 cells hidden

► ACCURACY MODEL

[] ↳ 1 cell hidden

► CLUSTERING MODEL

```
[ ] ↳ 12 cells hidden
```

▼ PERORMING CLASSIFICATION

```
iris = pd.read_csv("Iris dataset.csv")

from sklearn.model_selection import train_test_split

# Dropping the target and species since we only need the measurements
X = iris.drop(['Id','Species'], axis=1)

# converting into numpy array and assigning petal length and petal width
X = X.to_numpy()[:, (2,3)]
y = iris['Id']

# Splitting into train and test
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.5, random_state=42)

from sklearn.linear_model import LogisticRegression

log_reg = LogisticRegression()
log_reg.fit(X,y)

LogisticRegression()

training_prediction = log_reg.predict(X_train)
training_prediction

array([ 77, 110,  60,  23,  99, 101,  23,  23,  63, 130,  14, 145,  23,
        23, 115,  99, 119, 135, 108, 119,  60,  23,  23, 135, 119,  23,
        23,  23, 135, 110,  23, 110, 119,  23, 135, 120, 115, 135, 145,
        14, 110, 107, 145,  80,  60,  74,  23,  74,  74,  23,  68, 119,
       119,  23,  99, 110, 123,  14, 110,  23, 135, 119, 123,  68, 123,
        74,  74, 115, 137,  14,  60, 115,  23,  68, 119])

test_prediction = log_reg.predict(X_test)
test_prediction

array([135,  23, 119,  77, 135,  23,  65, 145,  77,  63, 145,  14,  23,
        14,  23, 147, 119,  68,  74, 110,  14, 115,  23, 119, 119, 145,
       123, 119,  23,  14,  23,  23,  74,  14,  23, 115,  77,  23,  23,
        14, 145, 122, 134,  23,  23,  63, 135, 123,  74, 119,  91, 119,
        77,  23, 119,  68,  23,  23,  14,  80, 115,  13,  14,  23,  91,
        23,  68, 119,  23,  74, 115,  23, 115, 145,  68])
```

```
from sklearn import metrics

print("Precision, Recall, Confusion matrix, in training\n")

# Precision Recall scores
print(metrics.classification_report(y_train, training_prediction, digits=3))

# Confusion matrix
print(metrics.confusion_matrix(y_train, training_prediction))
```

```
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Und
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Und
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Und
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Und
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Und
_warn_prf(average, modifier, msg_start, len(result))
```

```
print("Precision, Recall, Confusion matrix, in testing\n")
```

```
# Precision Recall scores
```

```
print(metrics.classification_report(y_test, test_prediction, digits=3))
```

```
# Confusion matrix
```

```
print(metrics.confusion_matrix(y_test, test_prediction))
```

macro avg	0.017	0.049	0.021	75
weighted avg	0.019	0.053	0.023	75

```
[[0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 0 0 0]
 ...
 [0 0 0 ... 0 0 0]
 [0 0 0 ... 1 0 0]
 [0 0 0 ... 0 0 0]]
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Und
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Und
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Und
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Und
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Und
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Und
_warn_prf(average, modifier, msg_start, len(result))
```

Conclusion: We have performed eda using different models where regression model have highest accuracy then a Other models